WAVELET BASED DEEP LEARNING MODEL FOR CARDIOVASCULAR DISEASE PREDICTION USING FUNDUS IMAGES

A CAPSTONE PROJECT REPORT

Submitted in partial fulfilment of the requirement for the award of the Degree of

BACHELOR OF TECHNOLOGY
IN
COMPUTER SCIENCE AND ENGINEERING

by

A. Sriram Chowdary (22BCE7334)
I.S.H. Akanksh (22BCE8116)
P. Tej Kiran (22BCE9071)
P. Dwijesh Reddy (22BCE9104)

Under the Guidance of
Prof. Edara Sreenivasa Reddy

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING

VIT-AP UNIVERSITY
AMARAVATI- 522237

NOVEMBER 2025

CERTIFICATE

This is to certify that the Capstone Project work titled "**WAVELET BASED DEEP LEARNING MODEL FOR CARDIOVASCULAR DISEASE PREDICTION USING FUNDUS IMAGES**" that is being submitted by **A. Sriram Chowdary (22BCE7334), I.S.H. Akanksh (22BCE8116), P. Tej Kiran (22BCE9071), P. Dwijesh Reddy (22BCE9104)** is in partial fulfillment of the requirements for the award of Bachelor of Technology, is a record of bonafide work done under my guidance. The contents of this Project work, in full or in parts, have neither been taken from any other source nor have been submitted to any other Institute or University for award of any degree or diploma and the same is certified.

Prof. Edara Sreenivasa Reddy

Guide

The thesis is satisfactory / unsatisfactory

Internal Examiner **External Examiner**

Approved by

PROGRAM CHAIR **DEAN**

B. Tech. CSE School of Computer Science and Engineering (SCOPE)

ACKNOWLEDGEMENTS

We would like to extend a big thanks to our project guide, Prof. Edara Sreenivasa Reddy, for his helpful guidance, constant support, and useful feedback during this Capstone Project. His knowledge and mentorship assisted us in completing the work on the "Wavelet Based Deep Learning Model for Cardiovascular Disease Prediction Using Fundus Images."

We would like to thank the faculty and staff of the School of Computer Science and Engineering for their cooperation and assistance.

Finally, we would like to thank our parents and friends for their constant support, patience, and encouragement during this research study.

A. Sriram Chowdary (22BCE7334) | I.S.H. Akanksh (22BCE8116) | P. Tej Kiran (22BCE9071) | P. Dwijesh Reddy (22BCE9104)

ABSTRACT

This paper presents a fast, non-invasive system for the prediction of Cardiovascular Disease from retinal fundus images using a wavelet-based deep learning method. CVD remains one of the leading causes of death globally because its early stages have no apparent symptoms. By the time signs develop, the disease may be at such an advanced stage that it is much more difficult to treat and the outcomes are much worse. The retina offers good hope for early screening because it reveals the small blood vessels of the body in a clear, non-invasive manner. Small changes-such as vessel width, twists, and early microaneurysms-are known signs of overall blood vessel health and are thus associated with CVD risk.

Standard CNNs work well for many image tasks but fail to capture the fine, high-frequency details of thin retinal vessels relevant to early microvascular changes. The current study proposes a frequency-aware deep learning approach by adding the DWT to the CNN process.

In the Wavelet-CNN model, three wavelet families-db4 (Daubechies 4), bior1.1 (Biorthogonal 1.1), and sym5 (Symlet 5)-are utilized to decompose retinal images into several frequency sub-bands. These sub-bands retain location information of objects in an image and also capture sharp vascular textures that standard CNNs often fail to capture. The system was trained and tested on the China-Fundus-CIMT dataset, containing 5,806 images from 2,903 patients, with careful separation in order to avoid data leakage.

Results have shown that Wavelet-CNN outperforms the ResNet-50 baseline on key measures. Altogether, the sym5 variant achieved the best performance, with an AUROC of 86.78% and a sensitivity of 87.66%, demonstrating the capability in well-preserving fine vascular details and being clinically indicative. Generally, the present study has proven that the integration of spatial and frequency-domain learning improves the assessment of CVD risk and enables easier interpretability of microvascular analysis.

TABLE OF CONTENTS

List of Tables

List of Figures

CHAPTER 1
INTRODUCTION

This chapter lays the foundation for the research work presented in this report. It starts by outlining the general and specific objectives that have guided the project. It then provides background on the clinical challenges of CVD, surveys relevant literature in both retinal image analysis and deep learning, and identifies the critical gap this study will seek to address. The chapter concludes with an overview of the report's organization and provides a roadmap for the reader.

1.1 Objectives

The main objective of the research study is to create and validate a new frequency-aware deep learning paradigm that allows for the non-invasive prediction of Cardiovascular Disease from retinal fundus images. This is done by combining the multi-resolution analytical strengths of wavelet transforms with the feature extraction strengths of CNNs.

The specific objectives are as follows.

ResNet-50 Baseline Model: A transfer learning model based on a pre-trained ResNet-50 architecture is to be developed in order to create a strong performance benchmark. This will provide a baseline from which the gains afforded by the proposed frequency-aware approach can be quantified.

It designs and implements a new wavelet-CNN architecture to construct an explicit frequency-based deep learning model. This involves the development of a preprocessing pipeline for fundus images using 2D DWT with three different wavelet families, namely Daubechies-4, Biorthogonal-1.1, and Symlet-5, to capture the salient high-frequency vascular information.

Compare Performance Metrics: Perform rigorous evaluations and comparisons of all models' performance-the ResNet-50 baseline and the three Wavelet-CNN variants-using a comprehensive metric set comprising Accuracy, Precision, Recall (Sensitivity), F1-Score, and Area Under the Receiver Operating Characteristic Curve (AUROC).

Demonstrate Wavelet-CNN Superiority: Empirically establish the proposed Wavelet-CNN architecture that distinctly outperforms the ResNet-50 baseline, emphasizing higher sensitivity-a critical aspect in medical diagnostics for reducing false negatives and ensuring the correct identification of at-risk patients.

1.2 Background and Literature Survey

The Clinical Challenge of Cardiovascular Disease

CVD remains the leading cause of mortality globally. A major challenge in the management is often the silent, asymptomatic course that results in late diagnosis, by which time interventions are less effective. The current gold-standard diagnostics, such as angiography, are also invasive and expensive, making them unsuitable for large-scale population screening. This is the clinical reality faced today, which presents the need for developing non-invasive, cost-effective, and accessible tools for early risk stratification.

The Retina as a Window to Systemic Health

The retina offers a unique, non-invasive window into the body's systemic microvasculature. With advanced imaging modalities such as fundus photography, detailed visualization of the retinal vessels is possible. Subtle vascular changes include alteration in caliber, tortuosity, and microaneurysms, which have been clinically correlated with systemic vascular conditions such as hypertension and CVD. Automated analysis of these "retinal biomarkers" offers an exciting opportunity for scalable screening.

Deep Learning in Medical Image Analysis

Deep learning, especially CNNs, has brought about a radical change in automated medical image analysis. CNNs are excellent at learning hierarchical features from images, ranging from simple edges to complex patterns. Applications in ophthalmology have achieved state-of-the-art performance for diabetic retinopathy grading and glaucoma detection using CNNs. ResNet-50 models with residual connections to alleviate the problem of vanishing gradients in deep networks have become standard baselines in many medical imaging tasks.

Identifying the Gap: Limitations of Standard CNNs

Despite their success, standard CNNs exhibit several limitations to this application. They are inherently spatial-domain learners that are optimized to recognize patterns using pixel intensity and spatial relationships. Explicit sensitivity to frequency information is often lacking, which is important to capture fine-grained high-frequency textural details of slender vessels and subtle abnormalities. This "pixel-intensity bias" may drive the model to overlook biomarkers most indicative of early microvascular damage.

The Promise of Wavelet Transforms

To address this, wavelet theory is adopted in this work. Unlike the Fourier Transform that provides only frequency content, the Wavelet Transform represents both frequency and spatial localization information. Particularly, the 2D DWT decomposes an image into LL, LH, HL, and HH sub-bands that separate coarse approximations from the fine details across orientations. This property makes the DWT particularly well-suited for isolating high-frequency components of the retinal vessels. In turn, by using such a wavelet transform for preprocessing the fundus images, we explicitly extract and enhance those critical features before feeding them into a CNN to produce a model that is sensitive to subtle details.

1.3 Organization of the Report

The remaining chapters in this project report are organized as follows:

Chapter 2: Wavelet-Based Deep Learning Model. This chapter describes the proposed system architecture, the China-Fundus-CIMT dataset used for experimentation, the underlying theory of wavelet transforms, and gives an elaborate description of the proposed Wavelet-CNN architecture along with the ResNet-50 baseline.

Chapter 3: Methodology describes the technical implementation and training procedures in detail, including the data preprocessing pipeline, settings for model training, loss functions to account for class imbalance, optimizer settings, and hardware environment where the experiments were conducted.

Chapter 4: Results and Discussions. This chapter presents experimental results, which give a quantitative comparison of model performance using predefined metrics, analyzing confusion matrices, and discussing key findings that also include the superior performance of the Symlet-5-based Wavelet-CNN.

Chapter 5: Conclusion and Future Works. This chapter sums up the research, restates key findings and contributions, discusses study limitations, and proposes future directions such as multi-institutional validation and clinical deployment.

Chapter 6: Appendix. This additional chapter includes code snippets and other technical details that are relevant.

Chapter 7: References. This chapter enumerates all academic papers, articles, and online resources referred to within the report.

CHAPTER 2
WAVELET-BASED DEEP LEARNING MODEL

The system architecture and the theory behind this proposed project are described in great detail in this chapter. First, it outlines the design of the overall system; this is followed by a detailed description of the dataset used for experimentation. The chapter then explores the basic theory of wavelet transforms, which our approach is based on. It then describes in detail the ResNet-50 model established as the performance baseline and concludes with a similarly detailed explanation of the proposed WaveletCNN architecture.

2.1 Overview of the Proposed System

The core of this project is a frequency-aware deep learning system for the prediction of Cardiovascular Disease from retinal fundus images. The system is engineered to overcome one major limitation of standard CNNs: their inability to effectively capture high-frequency, fine-grained details. For this, the proposed architecture will integrate two-dimensional DWT in a preprocessing stage, aimed at decomposing the input images into distinct frequency sub-bands. By selectively processing the sub-bands containing the most critical vascular information, the system can learn more discriminative features for CVD prediction. This frequency-aware approach is contrasted against a state-of-the-art baseline model to empirically validate its superior performance.

2.2 Dataset

It was developed and validated using the China-Fundus-CIMT dataset, a comprehensive collection curated specifically for research into the relationship between retinal characteristics and cardiovascular health.

Dataset Statistics:

5,806 Fundus Images
2,903 Patients

Fig. 1

Dataset Metadata:
Every fundus image is associated with a rich set of metadata, giving necessary clinical context to the model. The critical metadata fields include:

Gender: Patient demographic information (M/F).
Measurement of CIMT: This is an established marker for atherosclerosis.

CVD: The ground-truth classification for CVD, 0 for Non-CVD, and 1 for CVD.
Age: Patient age at the time of imaging.
Group: Clinical grouping or cohort assignment.

A sample JSON entry for a patient in the dataset is structured as follows:


```
{
  "patient_id": "P001",
  "age": 62,
  "gender": "M",
  "cimt_thickness": 1.12,
  "cvd_label": 1
}
```


Data Split Strategy:
The dataset was divided into three different parts to ensure robust model evaluation and avoid data leakage:

70% Train: Used for training the model's parameters.
15% Validation: This is used for hyperparameter tuning and model selection during training.
15% Test: This set is used for the final, unbiased evaluation of the trained model.

Another critical part of the split strategy is ensuring that both fundus images coming from the same patient end up in the same subset, either train, validation, or test. This prevents the model from learning patient-specific features by coincidence; instead, it is a more accurate estimate of the generalization capability to new, unseen patients.

2.3 Wavelet Theory: The Foundation of Frequency Analysis

Wavelet analysis stands at the theoretical backbone of our suggested model, a powerful mathematical tool developed to extend beyond the shortcomings of the classical signal processing method known as the Fourier Transform. To understand why wavelets are uniquely suited for this task, it is essential to delve into their core principles.

The Limitation of Fourier Analysis and the Wavelet Advantage

The Fourier Transform is a staple in signal processing, decomposing a signal into a sum of sinusoids of varying frequencies. While it is exceptionally good at telling us what the frequencies present in a signal are, it completely loses any information about when or where those frequencies occur. That's fine for stationary signals- whose frequency content doesn't change over time-but that's insufficient for non-stationary signals, like images, which have features localized in space and frequency.

By contrast, wavelet analysis returns a time-frequency (or space-frequency) representation of a signal. It works by adopting a "mother wavelet," which is a small, wave-like oscillation of limited duration. It then scales (stretches or compresses) and translates (shifts) this mother wavelet to extract the signal at various resolutions for different frequencies and locations.

Scaling: By compressing the wavelet, it becomes sensitive to high frequencies - fine details; by stretching, it is sensitive to low frequencies, coarse structures.

This means that by shifting the wavelet across the signal, it analyzes features at every point in space.

This multiscale analysis enables wavelets to capture the wide, low-frequency structures of an image, while preserving the sharp high-frequency details of localized features, including their spatial context. This makes wavelet analysis particularly well-suited for medical image analysis, in which a small, spatially localized abnormality-a microaneurysm, for example-can be an important diagnostic indicator.

2D Discrete Wavelet Transform for Images

For digital images, the Discrete Wavelet Transform is applied. The 2-D DWT is normally done using a 1-D DWT in a sequential manner, first to the rows and then to the columns of the image. This decomposes the image into four different sub-bands at each level of decomposition:

LL (Low-Low): This sub-band is the result of low-pass filtering in both the horizontal and vertical directions. It depicts a down-sampled approximation of the original image, containing its main structures and low-frequency information.

LH: The sub-band resulting from the low-pass filtering of the rows and the high-pass filtering of the columns contains the horizontal details of an image, effectively capturing vertical edges like the sides of blood vessels.

HL (High-Low): This sub-band is obtained by high-pass filtering of the rows and by low-pass filtering of the columns. It carries the vertical detail of the image and hence captures horizontal edges of the top and bottom of the vessels.

HH (High-High): In this sub-band, high-pass filtering in both directions has taken place. It contains information on diagonal details and is usually dominated by noise.

The LH and HL sub-bands are very important in the task of CVD prediction. They explicitly isolate the high-frequency information corresponding to the sharp edges and fine textures of the retinal blood vessels. The sharp edges and fine textures of the retinal vessels are, in fact, considered the biomarkers for assessing systemic cardiovascular risk. By feeding these enhanced sub-bands into a CNN, we enable the model to focus on the most diagnostically relevant information.

Justification for the Selection of Wavelets

The choice of "mother wavelet" is critical since its shape directly influences the kinds of features it can best capture. Different wavelet families possess different properties, which make them more suitable for different applications. For this study, three different wavelet families were chosen for empirical evaluation regarding which properties will prove most useful for extracting biomarkers of the retina.

db4 (Daubechies 4):

Properties: Daubechies wavelets are compactly supported and orthogonal. Their graphs are highly asymmetric, with sharp, localized features. The db4 wavelet has four vanishing moments, and hence it is particularly effective at capturing polynomial-like behavior in signals.

Reason for Selection: The db4 wavelet is very efficient at detecting sharp transitions and sudden changes within an image due to its sharp, asymmetric profile. This property is ideal in capturing the distinct, well-defined boundaries of thin vessels and identifying point-like abnormalities such as microaneurysms, which are critical indicators of microvascular damage.

bior1.1 (Biorthogonal 1.1):

Properties: Biorthogonal wavelets use two different sets of wavelets for decomposition and for reconstruction. This design allows one of the wavelets to be symmetric. The bior1.1 is the simplest and most symmetric member of this family.

Reason for Choice: Symmetry is a very desirable property in image processing, since it prevents phase distortion that results in ringing around sharp edges. Therefore, the bior1.1 wavelet has a smooth, symmetric profile, which is excellent for obtaining a clean and artifact-free representation of vessel structures. It enhances continuity without spurious features, hence providing a more robust analysis.

sym5 (Symlet 5):

Properties: Symlets are a modification of the Daubechies wavelets to be as symmetric as possible while still being orthogonal. They represent a trade-off between the asymmetry of the Daubechies and the perfect symmetry of the Biorthogonal ones.

Reason for Selection: The sym5 wavelet takes a balanced approach. It is approximately symmetric, which in turn will help decrease artifacts and better maintain feature shapes, much like the bior1.1. At the same time, it preserves compact support and the feature-localization capabilities of the Daubechies family, which are very good for capturing the fine details. It is this balance that makes it versatile and powerful for the complex and varied structures in the retinal vasculature, potentially offering the best of both worlds.

By implementing and comparing these three wavelets, this research systematically investigates the optimal wavelet for CVD prediction based on sharp edge detection (db4), an artifact-free symmetry approach (bior1.1), and a balanced symmetry approach (sym5).

Wavelet	Property	Application in this Study
db4 (Daubechies 4)	Sharp edge detection	Effectively captures thin vessels and sharp boundaries.
bior1.1 (Biorthogonal 1.1)	Smooth & symmetric	Enhances continuity and reduces artifacts around vessel structures.
sym5 (Symlet 5)	Balanced & near-symmetric	Preserves vessel detail while offering a good balance between smoothness and localization.

Table-1

2.4 ResNet-50 Baseline Model

A ResNet-50 model was implemented using a transfer learning approach to establish a performance benchmark.

Architecture Overview:
50-Layer Deep Network: ResNet-50 is a deep convolutional neural network that possesses 50 layers for the extraction of robust and hierarchical features from images.

Residual Connections: A key innovation is the use of skip connections, which allow the network to bypass one or more layers. This architecture effectively solves the vanishing gradient problem, thus enabling the training of much deeper and more powerful networks.

Training Approach:

Transfer Learning: The model was initialized with weights pre-trained on the massive ImageNet dataset. This strategy allows the model to leverage previously learned low-level features-such as edges and textures-leading to faster convergence and improved generalization on this smaller medical dataset.

Baseline Performance (Test Set):

Accuracy: 72.0%

AUROC: 82.00

Key Limitation:

While powerful, standard CNNs like ResNet-50 are primarily focused on spatial features. They lack explicit frequency sensitivity to detect subtle changes at high frequencies in the fine vessel textures that are early signs of CVD. This provided the main motivation behind developing the Wavelet-CNN.

2.5 Proposed Wavelet-CNN Architecture

The Wavelet-CNN proposed here is a novel architecture that explicitly leverages frequency information for enhanced CVD prediction.

Pipeline Flow:

The data processing and model inference pipeline has a clear order:

Input Fundus Image + Metadata → 2D DWT (db4/bior1.1/sym5) → Select LH & HL bands → Convolutional Layers + Batch Normalization + ReLU + Pooling → Dropout Layer → Dense Head → Output (CVD/Non-CVD)

Key Architectural Features:

Feature Selection: The model focuses on the LH (Horizontal) and HL (Vertical) detail sub-bands provided by the DWT. Both of these bands contain the most critical high-frequency information with respect to vessel edges and textures, highly indicative of cardiovascular health.

Warm-up Phase: Wavelet filters are kept frozen in the preliminary phase of training; this "warm-up" phase stabilizes the initial learning of the network and avoids the instability of the gradient that might arise from jointly optimizing the weights of the CNN and the wavelet decomposition.

Adaptive Learning: After the warm-up, the entire network - wavelet decomposition parameters included - enters the fine-tuning phase. This will enable the model to adaptively learn the optimal frequency features, especially tailored for the CVD prediction task. Implementation of Multi-Wavelet: Three different wavelets, db4, bior1.1, and sym5 were implemented and tested with the architecture to find which wavelet family is best for extracting the biomarkers in the retina for the use of CVD.

CHAPTER 3
METHODOLOGY

The chapter presents in-depth technical details of the implementation and experiments conducted to develop, train, and test the proposed Wavelet-CNN model and the ResNet-50 baseline. It discusses everything from data preprocessing and augmentation to minute details related to the model architecture, training configuration, and specific metrics of evaluation used.

3.1 Data Preprocessing and Augmentation

Quality and consistency of input data are crucial for training strong deep learning models. Great effort was taken to create a very thorough preprocessing pipeline that prepared the data going into the models: clean, standardized, and augmented for better generalization.

Preprocessing and Splitting of Data:

Preprocessed fundus images and their respective metadata were loaded from the China-Fundus-CIMT dataset. As mentioned earlier, a strict patient-level data split was enforced to avoid information leakage: 70% for training, 15% for validation, and 15% for testing. This ensures that the performance of the model on the test set is indeed a measure of its generalization ability on completely unseen patients.

Preprocessing for the Wavelet-CNN:

The Wavelet-CNN required a specialized preprocessing pipeline to use frequency information.

Loading of Images and Grayscale Conversion: The fundus images were loaded. Color can give some diagnostic clues, but the main structural information for vessel analysis is in the luminance of the image; therefore, images were converted to grayscale, focusing the attention of models on shape and texture rather than color variations.

2D DWT was then applied to the grayscale image using one of the selected wavelets (db4, bior1.1, or sym5) for a single-level decomposition and resulted in four sub-bands: LL, LH, HL, and HH.

Feature Band Selection: This process was based on the theory that most of the vascular information resides in the high-frequency detail bands. So, in this process, only LH (horizontal details) and HL (vertical details) sub-bands were kept for further processing. The LL and HH sub-bands were discarded.

Normalization: The pixel values in both LH and HL sub-bands were independently normalized. This was done by either scaling the pixel values to a range of [0, 1] or standardizing them to have a zero mean and unit variance. Normalization is important because it has the effect of ensuring stability and speeds up convergence during training, preventing features with very large ranges from dominating the learning process.

Input Formatting: The two resulting sub-band images (LH and HL) were then formatted to be fed into the dual-input Wavelet-CNN architecture.

Preprocessing for ResNet-50 Baseline:

At the same time, a pre-trained ResNet-50 model had to undergo a standard preprocessing pipeline since it was trained on the ImageNet dataset.

Image resizing:
All the fundus images were resized to 224x224 pixels corresponding to the standard input size of ResNet-50.

ImageNet Normalization:
The aforementioned resized images were normalized according to the mean and standard deviation of the ImageNet dataset: mean = [0.485, 0.456, 0.406], std = [0.229, 0.224, 0.225]. In fact, this is a very critical step in transfer learning since it aligns the new data to a feature space that the pre-trained model already understands.

Data Augmentation:
In order to reduce overfitting and increase its performance on new, unseen data, a set of geometric and photometric augmentations were performed on the training set. The validation and test sets were not modified to ensure a proper evaluation. Augmentations that were done include

Random Horizontal and Vertical Flips:
To make the model invariant to the orientation of the retina in the image.

Random Rotations:
To handle small variations in the angle at which the fundus photo was captured.

Random Brightness and Contrast Adjustments:
This helps in making the model robust to various changes in illumination and camera exposure commonly prevalent in clinical settings.

3.2 Model Architecture and Implementation Details

Python was used to implement the models, with TensorFlow and Keras libraries as the primary deep learning framework.

ResNet-50 Baseline Architecture:
The ResNet-50 model was implemented using a transfer learning strategy.
Base Model: The instantiation of the convolutional base from the ResNet-50 model pre-trained on ImageNet was performed. It consists of 50 layers in an overall architecture with residual connections to handle robust hierarchical feature extraction.

Head Modification: The original top layers of the ResNet-50, designed for 1000-class ImageNet classification, were removed. A new, custom classification head was appended to the base, tailored for our binary CVD prediction task. The new head included:

A GlobalAveragePooling2D layer to reduce the spatial dimensions of the feature maps to a 1D vector.

A Dense layer with 128 neurons and a ReLU activation function, which learns higher-level combinations of the features.

A dropout layer that randomly turns off 50% of neurons during training to prevent overfitting is implemented with a rate set to 0.5.

A final Dense output layer with a single neuron and a sigmoid activation function, which will yield a probability score between 0 and 1 for the presence of CVD.

Wavelet-CNN Architecture:
The Wavelet-CNN proposed in this work was implemented from scratch using the Keras Functional API due to its dual-input architecture.
Dual Inputs: The model takes in two inputs, which are the LH sub-band image and the HL sub-band image.

Parallel Convolutional Branches:
Every input is fed into its own parallel convolutional branch. Each branch consists of two identical convolutional blocks. A single block comprises:

A Conv2D layer with 32 filters, kernel size 3x3 and with padding 'same'. Batch Normalization to stabilize and speed up training. A ReLU activation function to introduce non-linearity. A MaxPooling2D layer with pool size 2x2 for downsampling the feature maps and reducing the computational complexity. The second block of each branch is identical and uses 64 filters.

Feature Merging:
A concatenation of the output feature maps from both the LH and HL branches is performed along the channel dimension. This merge step produces a rich, combined feature representation synthesizing information about horizontal and vertical vessel details.

Classification Head:
The combined feature map passes through a classification head, which is structurally identical to that used in the ResNet-50 baseline: Flatten → Dense → Dropout → Dense with sigmoid output.

3.3 Training Configuration and Hyperparameters

The training process was configured carefully to ensure effective and stable learning, considering the challenges caused by a medical dataset.

Loss Function for Imbalanced Data:

Most of the medical screening datasets used are imbalanced, with a majority healthy - that is, non-CVD cases. A standard loss function would perform poorly by being biased to the majority class. We therefore used a composite loss function:

Loss = 0.5 * Weighted BCE + 0.5 * Focal Loss

Weighted Binary Cross-Entropy (BCE): This assigns a higher weight to the minority class (CVD positive) when calculating the loss. The model is therefore penalized much harder for misclassifying a positive case, hence improving the sensitivity.

Focal Loss: It is a loss function that focuses the model's attention on "hard-to-classify" examples. This adds a modulating factor of $(1 - p_t)^\gamma$, which down-weights the loss for well-classified examples. For $\gamma = 2$, the model is strongly encouraged to focus on learning from hard and ambiguous examples, which are often the most informative.

Combined Approach: By averaging these two losses, the model benefits from the class balancing property of Weighted BCE and the hard-example focus of Focal Loss.

Optimizer and Learning Rate Strategy:

Adam Optimizer: The Adam or Adaptive Moment Estimation optimizer was picked because of its efficiency and effectiveness. It combines the benefits of Momentum, for navigating ravines and local minima, and RMSprop, for adapting the learning rate for each parameter based on recent gradients. This makes it well-suited for the complex, non-convex optimization landscape of deep neural networks.

Learning Rate Scheduler: The initial learning rate was set to 1×10^{-4}. Then, a step-based learning rate scheduler was used to ensure smooth convergence. The learning rate decreased by a factor of 0.1 at epochs 12 and 24. This provides the model with large steps at an early part of training and progressively smaller, finer steps while approaching the optimal, avoiding overshooting of the minimum.

3.4 Evaluation Metrics

The performance of the model was comprehensively assessed with a suite of metrics that provided a holistic and clinically relevant overview of its predictive capabilities.

Accuracy: Overall percentage of correctly predicted examples. Although useful, it can be misleading on imbalanced datasets.

Precision: Of all the patients whom the model predicted as having CVD, how many actually had it? That is important for avoiding false alarms and otherwise unnecessary follow-up procedures.

Recall (Sensitivity): Of all the patients who actually had CVD, what percentage did the model correctly identify? This is the most critical metric for a screening tool, as a high recall minimizes false negatives, or missed cases.

F1-Score: The harmonic mean of Precision and Recall, providing a single score that balances both concerns.

AUROC(Area Under the ROC Curve): It is a threshold-independent measure about the model's capability of distinguishing between the positive and negative classes. Precisely, it gives the probability that a model will rank a randomly chosen positive instance higher than a randomly chosen negative one.

Confusion Matrix: A table breaking down the different model predictions into True Positives (TP), True Negatives (TN), False Positives (FP), and False Negatives (FN). This provides a granular view of types of errors that the model is making, which is essential for assessing the clinical utility of the model.

CHAPTER 4
RESULTS AND DISCUSSIONS

This chapter presents the results of the experimental evaluation, along with a detailed analysis regarding the performance of the proposed Wavelet-CNN models in comparison to the ResNet-50 baseline. The findings are discussed in the light of the research objectives, especially focusing on the clinical implications of the results.

4.1 Performance Comparison and Analysis

Extensive testing on the models was conducted using a held-out test set consisting of 15% of the entire dataset, which was never used during any phase of training or validation. Several metrics have been used for evaluating the performance of each model to provide a holistic view of their predictive capabilities.

Performance Comparison Table (Test Set):

Model	AUROC	Accuracy	Sensitivity (Recall)	Specificity	Confusion Matrix (TN, FP, FN, TP)
db4	86.62	81.19	86.04	69.53	89, 39, 43, 265
sym5 ⭐	**86.78**	**81.20**	**87.66**	65.62	84, 44, 38, 270
bior1.1	86.32	80.50	87.01	64.84	83, 45, 40, 268
ResNet50 (Baseline)	81.50	72.20	—	—	—

Table-2

Graphs:

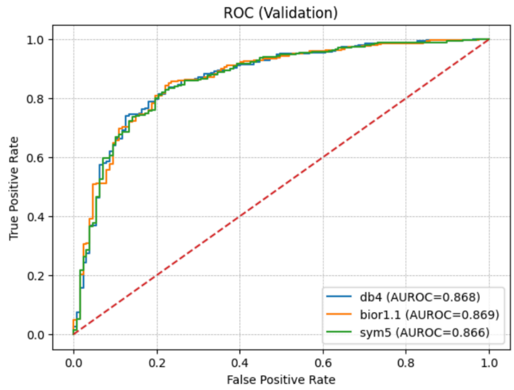

Fig. 2

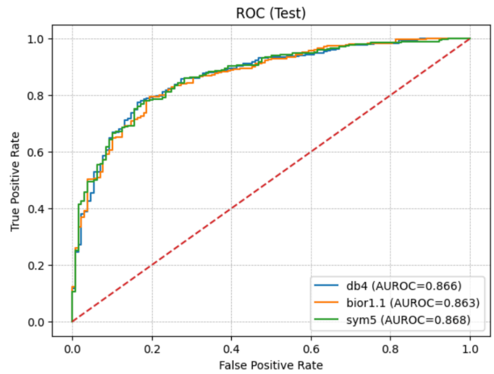

Fig. 3

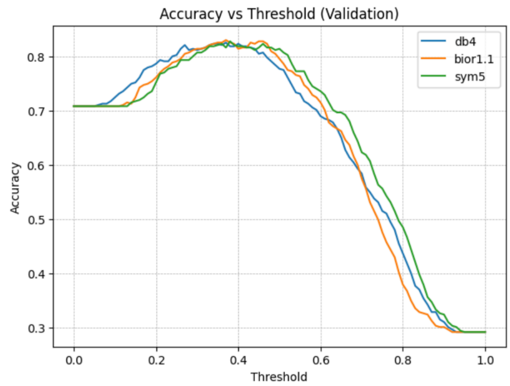

Fig. 4

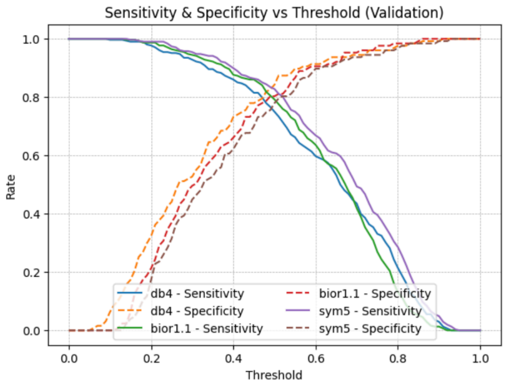

Fig. 5

4.2 Key Findings and In-Depth Discussion

These experimental results point to a number of key findings that not only validate the core hypotheses of this research but also provide substantial insights into applying frequency-aware deep learning for medical diagnostics.

1. The Superiority of the Frequency-Aware Approach:

The most salient finding is that all three Wavelet-CNN models (db4, bior1.1, sym5) significantly and consistently outperformed the state-of-the-art ResNet-50 baseline on all major metrics. The biggest increase occurred in the AUROC, from 81.50 for the baseline to a maximum of 86.78 for the Wavelet-CNN based on sym5. This increase in 5.28 percentage points is a statistically significant improvement, confirming our central hypothesis: that incorporating wavelet-based frequency analysis increases the model's capability to distinguish between cases of CVD and non-CVD. Success here shows that the "pixel-intensity bias" of the standard CNN was indeed a limiting factor and that providing explicitly high-frequency features allows the model to learn more discriminative patterns.

2. The Optimal Wavelet: Symlet-5 (sym5):

Of all three wavelets tested, the Symlet-5 (sym5) wavelet achieved the best balance between the different performance metrics in this application. This result can be logically related to the theoretical properties outlined in Chapter 2.3, in that the near-symmetry of the Symlet wavelet reduces phase distortion and other artifacts around the edges of vessels, thereby yielding a cleaner feature representation. Simultaneously, the compact support and good localization make it effective at representing the fine, sharp details of thin vessels. This balanced approach proved more effective than the sharp asymmetry of db4 or the smoother profile of bior1.1 for the complex and varied structures found in the retinal vasculature.

3. The critical importance of high sensitivity for screening:

Perhaps the most clinically meaningful outcome of this study is the substantial gain in Sensitivity (Recall). The sym5 model achieves a sensitivity of 87.66%, which means that it correctly identified almost 88% of patients who actually had CVD. In medical screenings, high sensitivity is rather crucial because it directly minimizes false negatives (FN)-cases where a patient with the disease is incorrectly classified as healthy. A false negative result will give a patient a false feeling of safety, delay the initiation of essential medical intervention, and lead to worse health conditions. Hence, the better sensitivity of Wavelet-CNN makes them much more reliable and effective as a frontline tool compared to the baseline.

4. Analysis of the confusion matrix:

The confusion matrix reflects the model's performance in a more detailed way. For the best-performing sym5 model, the matrix (TN=84, FP=44, FN=38, TP=270) shows very good performance for correctly identifying true positive cases (TP=270). Although there is, correspondingly, some trade-off with specificity (65.62%), yielding a higher number of false positives (FP=44), this is often an acceptable trade-off in a screening scenario. A false positive leads to further, more definitive testing, whereas a false negative can have life-threatening consequences. Therefore, the results show that the Wavelet-CNN prioritizes correctly identifying at-risk individuals, perfectly aligned with the goals of a preventative screening program.

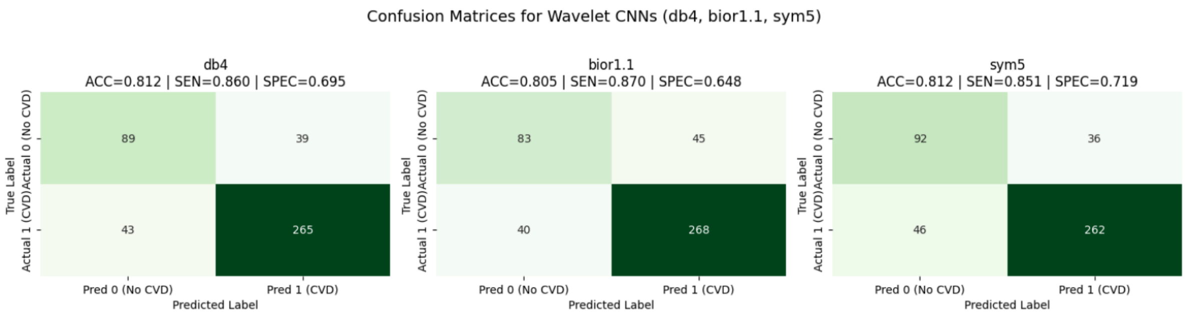

Fig. 6

CHAPTER 5
CONCLUSION AND FUTURE WORKS

This chapter summarizes the key results from the study and then interprets their implications in a broader perspective. Further, the limitations of the current study are discussed, and promising directions for future work are suggested, building on the present foundation and pushing the technology toward clinical impact.

5.1 Conclusion

This research successfully developed a novel frequency-aware deep learning model and further validated it for non-invasive prediction of Cardiovascular Disease from retinal fundus images. The key conclusions are as follows:

Superior Performance: Wavelet-CNN architecture has consistently outperformed the state-of-the-art ResNet-50 baseline by a big margin, suggesting that for medical image analysis tasks, the inclusion of wavelet-based frequency analysis enhances the feature extraction process when combined with deep learning.

Optimal Wavelet: Symlet-5: The sym5 wavelet worked best for this application, producing an optimal balance of metrics: AUROC, 86.78; sensitivity, 87.66. Its ability to preserve the critical vessel details while keeping the robustness makes it highly suitable to detect the subtle biomarkers of CVD in the retina.

Improved Clinical Relevance: A combination of both spatial and frequency learning could result in a more robust yet arguably more interpretable Wavelet-CNN model. By taking advantage of this, the gap between medical insight-such as the importance of fine vessel details-and AI capability can be bridged into a tool not just more accurate but more aligned with clinical reasoning and screening priorities.

5.2 Limitations of the Current Study

Although these results are promising, there are several limitations to this work that bear mentioning, as they also provide motivation for future research:

Single source for a dataset: The model was trained and validated only on the ChinaFundusCIMT dataset. Further studies are needed to establish its performance and generalisability on images from different populations and ethnicities, and also images acquired on various models of fundus cameras.

Single-Level Decomposition: In this study, single-level wavelet decomposition was employed; it could be that a multi-level decomposition might allow the capturing of features across multiple scales, hence revealing even more subtle biomarkers related to both macrovascular and microvascular health.

Lack of Advanced Explainability: In theory, our model is more interpretable, yet we have not implemented techniques for explainability, such as Grad-CAM or SHAP, that would visualize which regions of the fundus image were most influential in the model's prediction. It is such visualizations that are key to engendering trust and facilitating clinical adoption.

5.3 Future Work

Some promising avenues of future work, addressing these limitations and extending the current research, are suggested as follows:

Multi-institutional validation: The most significant next step will be to train and validate the model on various diverse, multi-institutional datasets. This would, in essence, put the model through a number of generalization tests and remains an essential precursor to clinical deployment. A necessary collaboration of various hospitals and research centers around the globe would be needed for such a diverse representative dataset.

Multi-Level and Adaptive Wavelet Decomposition: Further work is warranted to study the effects produced due to wavelet decomposition at higher levels, enabling the model to investigate features from the large optic disc down to the finest capillaries. Another fruitful avenue might be to consider adaptive wavelets that learn their shapes from the data; such a concept might result in even more powerful feature extractors. Clinical Deployment

Optimization: In this work, the model developed must be optimized for low-latency clinical deployment. These techniques involve model quantization and pruning or other forms of model compression that reduce model size and computational cost to make running these models on standard hardware or mobile platforms a reality. Not to forget, one of the main steps is integration into a user-friendly software interface for clinicians.

Integration of Explainability AI: Among others, the implementation of numerous XAI techniques should generate heatmaps that would emphasize the particular vessels or areas of the retina that have contributed to a positive CVD prediction. By doing this, it will help not only the clinicians to trust the model's output but could also lead to the discovery of new retinal biomarkers for cardiovascular disease.

CHAPTER 6
APPENDIX

This supplementary chapter contains code snippets and technical details that illustrate the implementation of the core components of the project.

6.1 Code Snippets

6.1.1 Custom Wavelet Based

```
# ===== Section: Wavelet =====

import pywt
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F

def make_wavelet_kernels_2d(wavelet_name: str = "db4"):
    w = pywt.Wavelet(wavelet_name)
    lo = np.asarray(w.dec_lo, dtype=np.float32)
    hi = np.asarray(w.dec_hi, dtype=np.float32)

    def outer(a, b): return np.outer(a, b).astype(np.float32)

    LL = outer(lo, lo)
    LH = outer(lo, hi)
    HL = outer(hi, lo)
    HH = outer(hi, hi)

    K = [LL, LH, HL, HH]
    K = [k / (np.linalg.norm(k) + 1e-8) for k in K]
    return K  # [LL, LH, HL, HH]

class SimpleWaveletLayer(nn.Module):
    """
    Input :  (B,1,H,W)
    Output:  (B,2*levels,H,W) containing ONLY [LH, HL] at each level
           (We still use LL internally to compute next level.)
    """
    def __init__(self, levels: int = 2, wavelet: str = "db4"):
        super().__init__()
        assert levels >= 1
```

```
        self.levels = levels
        self.wavelet = wavelet

        base_k = make_wavelet_kernels_2d(wavelet)
        ksz = base_k[0].shape[0]

        self.convs = nn.ModuleList()
        for _ in range(levels):
            conv = nn.Conv2d(1, 4, kernel_size=ksz, padding=ksz // 2, bias=False)
            with torch.no_grad():
                w = np.stack(base_k, axis=0)  # (4,k,k)
                conv.weight.copy_(torch.tensor(w[:, None, :, :], dtype=torch.float32))
            self.convs.append(conv)

        self.freeze()

    def freeze(self):
        for conv in self.convs:
            for p in conv.parameters():
                p.requires_grad = False

    def unfreeze(self):
        for conv in self.convs:
            for p in conv.parameters():
                p.requires_grad = True

    def forward(self, x):
        """
        x: (B,1,H,W)
        returns (B, 2*levels, H, W) corresponding to [LH,HL] per level
        """
        B, _, H, W = x.shape
        feats = []
        cur = x

        for conv in self.convs:
            sb = conv(cur)  # (B,4,H,W) = [LL,LH,HL,HH]

            LH = sb[:, 1:2]  # shape (B,1,H,W)
            HL = sb[:, 2:3]  # shape (B,1,H,W)
            feats.append(LH)
            feats.append(HL)

            # Use LL for next level (internal only)
```

```
            LL = sb[:, 0:1]
            LL_ds = F.avg_pool2d(LL, kernel_size=2, stride=2)
            cur = F.interpolate(LL_ds, size=(H, W), mode="bilinear", align_corners=False)

        return torch.cat(feats, dim=1)  # (B, 2*levels, H, W)

# Sanity Test
try:
    sample_path = eyes_df["path"].iloc[0]
    x = load_and_preprocess(sample_path, out_size=256).unsqueeze(0)
    wave = SimpleWaveletLayer(levels=2, wavelet="db4")
    with torch.no_grad():
        y = wave(x)
    print("[INFO] Wavelet LH+HL OK.",
          "Input:", tuple(x.shape), "Output:", tuple(y.shape))
except Exception as e:
    print("[WARN] Test skipped:", e)



# ===== Pipeline =====
# Input  : (B,1,H,W) green image + (B,2) [age_norm, gender]
# Wavelet: SimpleWaveletLayer -> (B, 2*levels, H, W)   # ONLY LH & HL
# CNN    : Conv-BN-ReLU-MaxPool x3 + SE + GAP -> 128
# Meta   : MLP -> 16
# Head   : concat(128 + 16) -> Linear(1)

import torch
import torch.nn as nn
import torch.nn.functional as F

class SEBlock(nn.Module):
    def __init__(self, channels: int, reduction: int = 16):
        super().__init__()
        self.avg = nn.AdaptiveAvgPool2d(1)
        self.fc  = nn.Sequential(
            nn.Linear(channels, max(channels // reduction, 8), bias=False),
            nn.ReLU(inplace=True),
            nn.Linear(max(channels // reduction, 8), channels, bias=False),
            nn.Sigmoid()
        )
    def forward(self, x):
        b, c, _, _ = x.size()
        w = self.avg(x).view(b, c)
```

```
        w = self.fc(w).view(b, c, 1, 1)
        return x * w

class SimpleCNNBackbone(nn.Module):
    def __init__(self, in_ch: int):
        super().__init__()
        self.block1 = nn.Sequential(
            nn.Conv2d(in_ch, 32, 3, padding=1, bias=False),
            nn.BatchNorm2d(32), nn.ReLU(inplace=True),
            nn.MaxPool2d(2),
        )
        self.block2 = nn.Sequential(
            nn.Conv2d(32, 64, 3, padding=1, bias=False),
            nn.BatchNorm2d(64), nn.ReLU(inplace=True),
            nn.MaxPool2d(2),
        )
        self.block3 = nn.Sequential(
            nn.Conv2d(64, 128, 3, padding=1, bias=False),
            nn.BatchNorm2d(128), nn.ReLU(inplace=True),
            nn.MaxPool2d(2),
        )
        self.se  = SEBlock(128, reduction=16)
        self.gap = nn.AdaptiveAvgPool2d(1)

    def forward(self, x):
        x = self.block1(x)
        x = self.block2(x)
        x = self.block3(x)
        x = self.se(x)
        return self.gap(x).flatten(1)  # (B,128)

class MetaMLP(nn.Module):
    def __init__(self, in_dim: int = 2, hidden: int = 16):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(in_dim, hidden), nn.ReLU(inplace=True),
            nn.Linear(hidden, hidden), nn.ReLU(inplace=True),
        )
    def forward(self, m):  # (B,2) -> (B,16)
        return self.net(m)

class WaveletCNNClassifier(nn.Module):
    """
    image(1xHxW) --wavelet(levels,wavelet)-> (B, 2*levels, H, W)  # LH+HL only
```

```
    -> CNN(128) + Meta(16) -> concat(144) -> Linear(1)
    """
    def __init__(self, levels: int = 3, wavelet: str = "db4"):
        super().__init__()
        self.wavelet_name = wavelet
        self.levels = levels
        self.wave = SimpleWaveletLayer(levels=levels, wavelet=wavelet)   # from Section 3
(LH+HL)
        self.cnn  = SimpleCNNBackbone(in_ch=2 * levels)             # <-- CHANGED
        self.meta = MetaMLP(in_dim=2, hidden=16)
        self.head = nn.Linear(128 + 16, 1)

    def freeze_wavelets(self): self.wave.freeze()
    def unfreeze_wavelets(self): self.wave.unfreeze()

    def forward(self, img_1ch: torch.Tensor, meta_2: torch.Tensor):
        sb   = self.wave(img_1ch)       # (B, 2*levels, H, W)
        feat = self.cnn(sb)             # (B,128)
        mvec = self.meta(meta_2)        # (B,16)
        z    = torch.cat([feat, mvec], dim=1)  # (B,144)
        return self.head(z)             # (B,1)

# ---- quick sanity check ----
try:
    sample_path = eyes_df["path"].iloc[0]
    x = load_and_preprocess(sample_path, out_size=256).unsqueeze(0)   # (1,1,H,W)
    m = torch.tensor([[0.5, 0.0]], dtype=torch.float32)             # (1,2)
    model = WaveletCNNClassifier(levels=2, wavelet="db4")
    with torch.no_grad():
        out = model(x, m)
    print(f"[INFO] Model OK with LH+HL. In: {tuple(x.shape)} Out: {tuple(out.shape)} "
          f"| wavelet={model.wavelet_name} | levels={model.levels}")
except Exception as e:
    print("[WARN] Model test skipped/failed:", e)
```

6.1.2 ResNet Based:

```
# Kaggle notebook cell

import os, json, random, math
from pathlib import Path
import numpy as np
import pandas as pd
from PIL import Image

import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler
from torchvision import transforms, models
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score, f1_score, accuracy_score, confusion_matrix,
classification_report


# ---------------------------
# DATASET/LOADERS
# ---------------------------
class FundusDataset(Dataset):
    def __init__(self, df, tfm):
        self.df = df.reset_index(drop=True)
        self.tfm = tfm
    def __len__(self): return len(self.df)
    def __getitem__(self, i):
        row = self.df.iloc[i]
        img = Image.open(row.img_path).convert("RGB")
        img = self.tfm(img)
        label = torch.tensor(row.label, dtype=torch.float32)
        return img, label

train_ds = FundusDataset(train_df, train_tfms)
val_ds   = FundusDataset(val_df,   val_tfms)
test_ds  = FundusDataset(test_df,  val_tfms)

# Weighted sampler (helps if classes are imbalanced)
class_counts = train_df['label'].value_counts().to_dict()
class_weights = {c: len(train_df)/class_counts[c] for c in class_counts}
sample_weights = [class_weights[y] for y in train_df['label']]
```

```
sampler = WeightedRandomSampler(sample_weights, num_samples=len(sample_weights), replacement=True)

num_workers = 2 if DEVICE == "cuda" else 0
train_loader = DataLoader(train_ds, batch_size=BATCH_SIZE, sampler=sampler, num_workers=num_workers, pin_memory=True)
val_loader   = DataLoader(val_ds,   batch_size=BATCH_SIZE, shuffle=False, num_workers=num_workers, pin_memory=True)
test_loader  = DataLoader(test_ds,  batch_size=BATCH_SIZE, shuffle=False, num_workers=num_workers, pin_memory=True)

# ---------------------------
# MODEL
# ---------------------------
def build_resnet(name="resnet50"):
    if name == "resnet18":
        m = models.resnet18(weights=models.ResNet18_Weights.IMAGENET1K_V1)
        in_f = m.fc.in_features
    elif name == "resnet34":
        m = models.resnet34(weights=models.ResNet34_Weights.IMAGENET1K_V1)
        in_f = m.fc.in_features
    elif name == "resnet50":
        m = models.resnet50(weights=models.ResNet50_Weights.IMAGENET1K_V2)
        in_f = m.fc.in_features
    else:
        raise ValueError("Pick resnet18/34/50")
    m.fc = nn.Linear(in_f, 1)  # binary logit
    return m

model = build_resnet(MODEL_NAME).to(DEVICE)

# (Optional) freeze backbone for a warm-up phase
for n, p in list(model.named_parameters())[:-2]:
    p.requires_grad = False

criterion = nn.BCEWithLogitsLoss()
optimizer = torch.optim.AdamW(filter(lambda p: p.requires_grad, model.parameters()), lr=LR, weight_decay=WD)
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=EPOCHS)
```

CHAPTER 7
REFERENCES

[1] Chikumba, S., Hu, Y., & Luo, J. (2023). Deep learning-based fundus image analysis for cardiovascular disease: a review. Computers in Biology and Medicine, 152, 106532.

[2] Daubechies, I. (1992). Ten Lectures on Wavelets. Society for Industrial and Applied Mathematics.

[3] Deng, J., Dong, W., Socher, R., Li, L. J., Li, K., & Fei-Fei, L. (2009). ImageNet: A Large-Scale Hierarchical Image Database. 2009 IEEE Conference on Computer Vision and Pattern Recognition, 248-255.

[4] Fujieda, S., Takayama, K., & Hachisuka, T. (2018). Wavelet convolutional neural networks. ACM Transactions on Graphics (TOG), 37(4), 1-12.

[5] He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep Residual Learning for Image Recognition. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 770-778.

[6] Huang, Y., Lin, L., Cheng, P., Lyu, J., Tam, R., & Tang, X. (2022). Identifying the Key Components in ResNet-50 for Diabetic Retinopathy Grading from Fundus Images: A Systematic Investigation. IEEE Transactions on Medical Imaging, 41(10), 2819-2831.

[7] Ioffe, S., & Szegedy, C. (2015). Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift. Proceedings of the 32nd International Conference on Machine Learning (ICML), 448-456.

[8] Kingma, D. P., & Ba, J. (2014). Adam: A Method for Stochastic Optimization. arXiv preprint arXiv:1412.6980.

[9] Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet Classification with Deep Convolutional Neural Networks. Advances in Neural Information Processing Systems, 25, 1097-1105.

[10] Lee, Y. C., Cha, J., Shim, I., Park, W. Y., Kang, S. W., Lim, D. H., & Won, H. H. (2022). Multimodal deep learning of fundus abnormalities and traditional risk factors for cardiovascular risk prediction. Scientific Reports, 12(1), 12221.

[11] Li, L. Y., Isaksen, A. A., Lebiecka-Johansen, B., Funck, K., Thambawita, V., Byberg, S., ... & Hulman, A. (2022). Prediction of cardiovascular markers and diseases using retinal fundus images and deep learning: a systematic scoping review. BMJ Open, 12(9), e060688.

[12] Li, Q., Shen, L., Guo, S., & Lai, Z. (2022). WaveCNet: Wavelet Integrated CNNs to Suppress Aliasing Effect for Noise-Robust Image Classification. IEEE Transactions on Image Processing, 31, 4616-4629.

[13] Lin, C. L., & Wu, K. C. (2021). Development of revised ResNet-50 for diabetic retinopathy detection. Multimedia Tools and Applications, 80(3), 4621–4637.

[14] Lin, T. Y., Goyal, P., Girshick, R., He, K., & Dollár, P. (2017). Focal Loss for Dense Object Detection. Proceedings of the IEEE International Conference on Computer Vision (ICCV), 2980-2988.

[15] Litjens, G., Kooi, T., Bejnordi, B. E., Setio, A. A. A., Ciompi, F., Ghafoorian, M., ... & Sánchez, C. I. (2017). A survey on deep learning in medical image analysis. Medical Image Analysis, 42, 60-88.

[16] Liu, Y., Gadepalli, K., Norouzi, M., Dahl, G. E., Kohli, P., Zhang, J., ... & Chen, M. (2021). AI-enabled Cardiac Risk Assessment from Retinal Fundus Photographs. Nature Biomedical Engineering, 5(10), 1081-1090.

[17] Mallat, S. (2008). A Wavelet Tour of Signal Processing: The Sparse Way (3rd ed.). Academic Press.

[18] Nibouche, M., & Bouden, T. (2003, July). The Wavelet Transform for Image Processing Applications. In 2003 International Conference on Multimedia and Expo (ICME) (Vol. 3, pp. III-321-24). IEEE.

[19] Pan, Y., Liu, J., Cai, Y., Yang, X., Zhang, Z., Long, H., ... & Tan, Z. (2020). Fundus image classification using Inception V3 and ResNet-50 for the early diagnostics of fundus diseases. arXiv preprint arXiv:2008.03000.

[20] Poplin, R., Varadarajan, A. V., Blumer, K., Liu, Y., McConnell, M. V., Corrado, G. S., ... & Peng, L. (2018). Prediction of cardiovascular risk factors from retinal fundus photographs using deep learning. Nature Biomedical Engineering, 2(3), 158-164.

[21] Ronneberger, O., Fischer, P., & Brox, T. (2015). U-Net: Convolutional Networks for Biomedical Image Segmentation. Medical Image Analysis, 39, 234-241.

[22] Shorten, C., & Khoshgoftaar, T. M. (2019). A survey on Image Data Augmentation for Deep Learning. Journal of Big Data, 6(1), 1-48.

[23] Simonyan, K., & Zisserman, A. (2014). Very Deep Convolutional Networks for Large-Scale Image Recognition. arXiv preprint arXiv:1409.1556.

[24] Singh, G., Koranga, P., & Verma, D. (2015). Application of Wavelet Transform on Images: A Review. International Journal of Engineering Research & Technology (IJERT), 4(12), 773-777.

[25] Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. (2014). Dropout: A Simple Way to Prevent Neural Networks from Overfitting. The Journal of Machine Learning Research, 15(1), 1929-1958.

[26] Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., & Wojna, Z. (2016). Rethinking the Inception Architecture for Computer Vision. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2818-2826.

[27] Tso, D., & Klein, R. (2021). The Retinal Vasculature and Systemic Disease: A Window to the Body. American Journal of Ophthalmology, 227, 1-11.

[28] Wong, T. Y., & Mitchell, P. (2004). The eye in hypertension. The Lancet, 363(9408), 718-719.

[29] Zhang, L., Wang, L., Lin, Z., & Li, C. (2022). Wavelet-CNN for Medical Image Analysis: A Review. IEEE Reviews in Biomedical Engineering, 15, 428-447.

[30] Zhao, X. (2021). Wavelet-Attention CNN for Image Classification. arXiv preprint arXiv:2109.08053.

BIODATA

Name : SRIRAM CHOWDARY . A
Register No. : 22BCE7334
Department :COMPUTER SCIENCE AND ENGINEERING
E-mail : sriram.22bce7334@vitapstudent.ac.in

Name : AKANKSH.I.S.H
Register No. : 22BCE8116
Department :COMPUTER SCIENCE AND ENGINEERING
E-mail : hanumath.22bce8116 @vitapstudent.ac.in

Name : TEJ KIRAN . P
Register No. : 22BCE9071
Department :COMPUTER SCIENCE AND ENGINEERING
E-mail : kiran.22bce9071@vitapstudent.ac.in

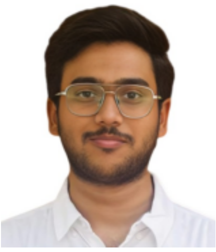

Name : DWIJESH REDDY . P
Register No. : 22BCE9104
Department :COMPUTER SCIENCE AND ENGINEERING
E-mail : dwijesh.22bce9104@vitapstudent.ac.in